

Pep/10

J. Stanley Warford

June 14, 2019

Here are the differences between Pep/10 and Pep/9 along with a rationale for each change.

1. STOP replaced by RET

The STOP instruction is no longer in the instruction set. Instead, the operating system now calls the C `main()` function with the system return value preset to 0. The translation more closely matches the terminating C statement

```
return 0;
```

The symbolic debugger of the Pep/10 IDE now shows the run-time stack from the OS call with two cells – `retAddr` and `retVal`. If students terminate their programs with RET the return value will be 0 because that is the preset return value, and control is returned to the simulator the same way a STOP instruction does in Pep/9. However, if they modify the value before the return, the OS issues an error message with an echo of the error number.

The operating system has a new dispatcher component as the interface between the OS and the application. This interface is more realistic of the way C works and reinforces the concept that the operating system calls the application, and the application returns control to the operating system.

2. Memory-mapped shutdown port

Pep/9 introduced the concept of memory-mapped I/O ports. In another step toward hardware realism, and to have a mechanism for terminating a simulation, Pep/10 has a memory-mapped shutdown port. If any value at all is written to the port the simulation is terminated and control is returned to the IDE.

Students first learn how to program in machine language at the ISA3 level without the assistance of the operating system. They learn how to store a byte to the output port with direct addressing to output an ASCII character. In Pep/10, they simply store a byte to the shutdown port with direct addressing to terminate their programs.

There are two benefits to this feature. First, students do not need to learn a new STOP instruction to terminate their machine language programs. But more importantly, they learn the utility of memory-mapped device registers with this rudimentary example.

3. Trap instructions replaced by system calls

Pep/9 has five trap instructions – NOP, DECI, DECO, HEXO, and STRO. Pep/10 replaces them with two system calls – SYCALL for system call and USYCALL for unary system call. For example, DECI is no longer an instruction mnemonic as it is in Pep/9, but a symbol exported from the operating system. This decimal input instruction in Pep/9

```
DECI    num,d
```

becomes the following system call in Pep/10

```
LDWT    DECI,i
SYCALL  num,d
```

Pep/10 has a new Trap register visible at level ISA3 used by the system call instructions. LDWT is the load word trap instruction, which loads the entry point address of the DECI code in the operating system.

In Pep/9, the use of the trap instructions so closely parallels the use of the native ISA instructions that beginning students frequently do not even realize they are system calls. Pep/10 has the pedagogic advantage of making system calls explicit, and is more realistic. From a system design perspective, the trap instructions no longer occupy the opcode space which opens the possibility for new instructions. The new design also does not constrain the number or variety of system calls.

4. Assembler macro facility

In yet another step toward industry standard practice, Pep/10 introduces assembler macro expansions. The IDE provides cononical standard macros for all the system calls. With the supplied macros, the above DECI example becomes simply

```
@DECI  num,d
```

where @DECI is now the macro name. The generated program listing shows the macro source instruction and its expansion. The convenient CHARI and CHARO instructions from Pep/8 and earlier are now back as macros @CHARI and @CHARO. This is a pedagogical improvement over Pep/9, because now character I/O is programmed exactly like decimal I/O. With dynamic allocation, students no longer need to copy/paste the code for `malloc()` at the end of their source because @malloc is a supplied macro.

It is possible for students to write their own macros with the Pep/10 IDE, which contains documentation for how to do so. However, writing macros is outside the scope of this text.

5. Easier modification of OS

In Pep/9, to write a new trap instruction the IDE requires you to redefine one of the existing mnemonics. Now that trap instructions are replaced with system calls there is no longer such a requirement. The Pep/10 assembler picks up all the information it needs from the OS source code using a new `.EXPORT` directive that makes OS symbols accessible to the application code. Students can write any number of system calls unconstrained by the opcode space.

6. New instruction `XORr`

Deletion of the trap instructions from Pep/9 opened up the opcode space for new instructions. Finally, the exclusive OR instruction `XORr` is a native ISA instruction.

7. Improved CPU

At the ISA level, the only difference in the CPU is the new 16-bit Trap register. At the LG1 level, the arithmetic/logic unit and the data paths for both the one-byte bus and the two-byte bus are unchanged from Pep/9. The register bank has an additional input port for a new D bus originating from the control section. Pep/10 has only two constant registers covering the most common cases instead of the awkward five in Pep/9. Now any constant value can be loaded into any register from the control section over the D bus.

8. Improved instruction set

The instruction set is now more representative of real ISA instruction sets. Neither the `STOP` instruction nor the specialized trap mnemonics of Pep/9 are in actual ISA sets. System calls are also typical. Both Pep/9 and Pep/10 have 40 ISA instructions, but the Pep/10 set is more regular. The opcode space is cleaner, with all the unary instructions followed by all the branch instructions with two addressing modes followed by the remaining instructions with eight addressing modes.

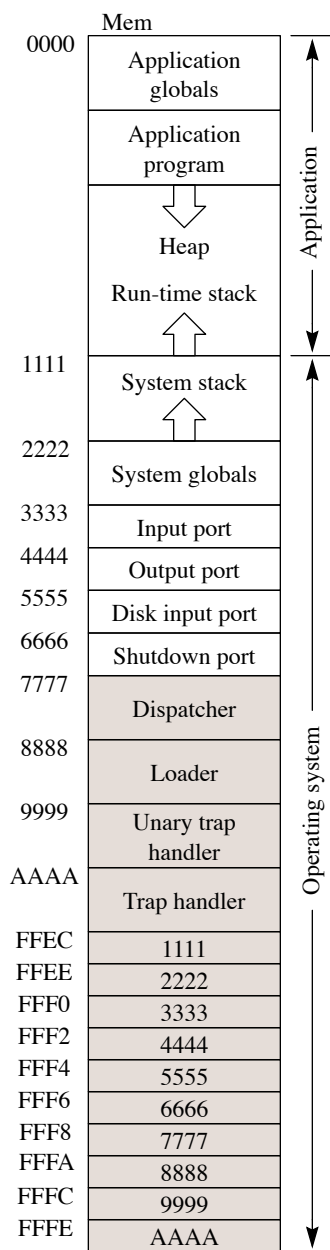
9. Improved control section

For the Pep/9 micro application, the microinstruction format now supports microcalls and microreturns made possible by a return address call stack in the control section.

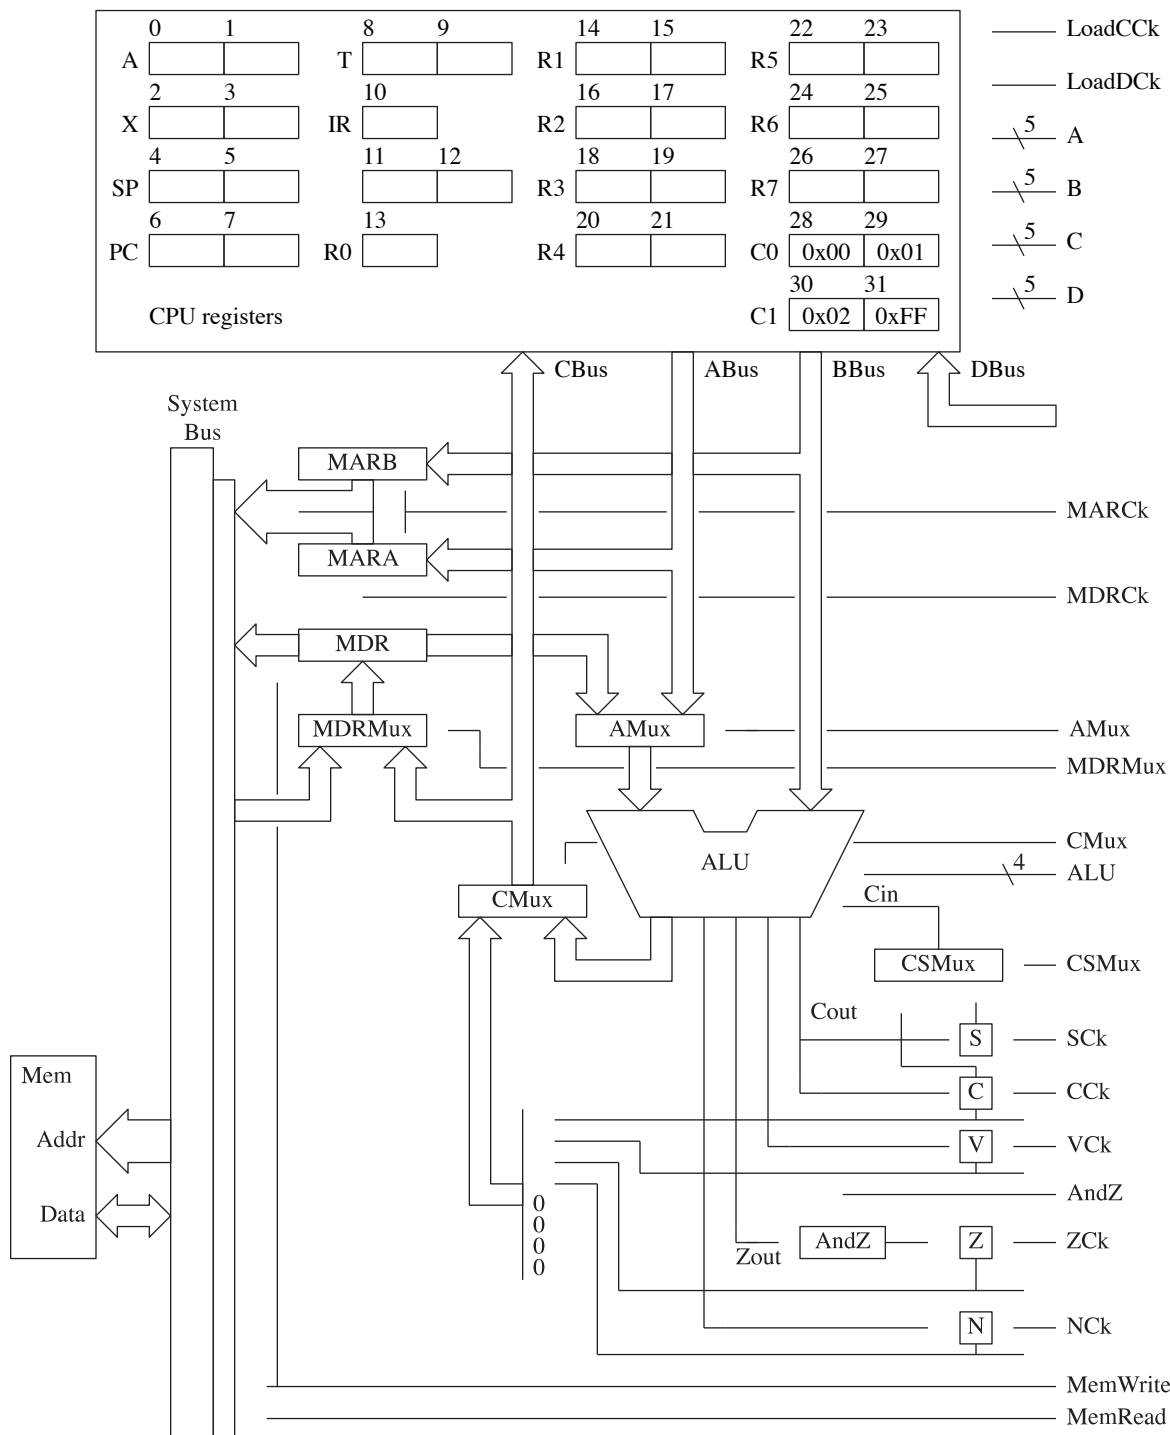
Instruction Specifier	Mnemonic	Instruction	Addressing Mode	Status Bits
0000 0000	RET	Return from CALL	U	
0000 0001	RETSY	Return from system CALL	U	
0000 0010	MOVSPA	Move SP to A	U	
0000 0011	MOVASP	Move A to SP	U	
0000 0100	MOVFLGA	Move NZVC flags to A(12..15)	U	
0000 0101	MOVAFLG	Move A(12..15) to NZVC flags	U	
0000 0110	MOVTPC	Move T to PC	U	
0000 0111	NOP	No operation	U	
0000 1000	USYCALL	Unary system call	U	
0001 000r	NOTr	Bitwise invert r	U	NZ
0001 001r	NEGr	Negate r	U	NZV
0001 010r	ASLr	Arithmetic shift left r	U	NZVC
0001 011r	ASRr	Arithmetic shift right r	U	NZC
0001 100r	ROLr	Rotate left r	U	C
0001 101r	RORr	Rotate right r	U	C
0001 110a	BR	Branch unconditional	i, x	
0001 111a	BRLE	Branch if less than or equal to	i, x	
0010 000a	BRLT	Branch if less than	i, x	
0010 001a	BREQ	Branch if equal to	i, x	
0010 010a	BRNE	Branch if not equal to	i, x	
0010 011a	BRGE	Branch if greater than or equal to	i, x	
0010 100a	BRGT	Branch if greater than	i, x	
0010 101a	BRV	Branch if V	i, x	
0010 110a	BRC	Branch if C	i, x	
0010 111a	CALL	Call subroutine	i, x	
0011 0aaa	SYCALL	System call	i, d, n, s, sf, x, sx, sfx	
0011 1aaa	LDWT	Load word T from memory	i, d, n, s, sf, x, sx, sfx	
0100 raaa	LDWr	Load word r from memory	i, d, n, s, sf, x, sx, sfx	NZ
0101 raaa	LDBr	Load byte r(8..15) from memory	i, d, n, s, sf, x, sx, sfx	NZ
0110 raaa	STWr	Store word r to memory	d, n, s, sf, x, sx, sfx	
0111 raaa	STBr	Store byte r(8..15) to memory	d, n, s, sf, x, sx, sfx	
1000 raaa	CPWr	Compare word to r	i, d, n, s, sf, x, sx, sfx	NZVC
1001 raaa	CPBr	Compare byte to r(8..15)	i, d, n, s, sf, x, sx, sfx	NZVC
1010 raaa	ADDr	Add to r	i, d, n, s, sf, x, sx, sfx	NZVC
1011 raaa	SUBr	Subtract from r	i, d, n, s, sf, x, sx, sfx	NZVC
1100 raaa	ANDr	Bitwise AND to r	i, d, n, s, sf, x, sx, sfx	NZ
1101 raaa	ORr	Bitwise OR to r	i, d, n, s, sf, x, sx, sfx	NZ
1110 raaa	XORr	Bitwise XOR to r	i, d, n, s, sf, x, sx, sfx	NZ
1111 0aaa	ADDSP	Add to SP	i, d, n, s, sf, x, sx, sfx	NZVC
1111 1aaa	SUBSP	Subtract from SP	i, d, n, s, sf, x, sx, sfx	NZVC

Instruction	Register transfer language specification
RET	$PC \leftarrow \text{Mem}[SP]; SP \leftarrow SP + 2$
RETSY	$NZVC \leftarrow \text{Mem}[SP]\langle 4..7 \rangle; A \leftarrow \text{Mem}[SP + 1]; X \leftarrow \text{Mem}[SP + 3]; PC \leftarrow \text{Mem}[SP + 5]; SP \leftarrow \text{Mem}[SP + 7]$
MOVSPA	$A \leftarrow SP$
MOVASP	$SP \leftarrow A$
MOVFLGA	$A\langle 8..11 \rangle \leftarrow 0, A\langle 12..15 \rangle \leftarrow NZVC$
MOVAFLG	$NZVC \leftarrow A\langle 12..15 \rangle$
MOVTPC	$PC \leftarrow T$
NOP	<i>{No operation}</i>
USYCALL	$Y \leftarrow \text{Mem}[\text{FFF0}]; \text{Mem}[Y - 2] \leftarrow SP; \text{Mem}[Y - 4] \leftarrow PC; \text{Mem}[Y - 6] \leftarrow X;$ $\text{Mem}[Y - 8] \leftarrow A; \text{Mem}[Y - 9]\langle 4..7 \rangle \leftarrow NZVC; SP \leftarrow Y - 9; PC \leftarrow \text{Mem}[\text{FFF8}]$
NOTr	$r \leftarrow \neg r; N \leftarrow r < 0, Z \leftarrow r = 0$
NEGr	$r \leftarrow -r; N \leftarrow r < 0, Z \leftarrow r = 0, V \leftarrow \{overflow\}$
ASLr	$C \leftarrow r\langle 0 \rangle, r\langle 0..14 \rangle \leftarrow r\langle 1..15 \rangle, r\langle 15 \rangle \leftarrow 0; N \leftarrow r < 0, Z \leftarrow r = 0, V \leftarrow \{overflow\}$
ASRr	$C \leftarrow r\langle 15 \rangle, r\langle 1..15 \rangle \leftarrow r\langle 0..14 \rangle; N \leftarrow r < 0, Z \leftarrow r = 0$
ROLr	$C \leftarrow r\langle 0 \rangle, r\langle 0..14 \rangle \leftarrow r\langle 1..15 \rangle, r\langle 15 \rangle \leftarrow C$
RORr	$C \leftarrow r\langle 15 \rangle, r\langle 1..15 \rangle \leftarrow r\langle 0..14 \rangle, r\langle 0 \rangle \leftarrow C$
BR	$PC \leftarrow \text{Oprnd}$
BRLE	$N = 1 \vee Z = 1 \Rightarrow PC \leftarrow \text{Oprnd}$
BRLT	$N = 1 \Rightarrow PC \leftarrow \text{Oprnd}$
BREQ	$Z = 1 \Rightarrow PC \leftarrow \text{Oprnd}$
BRNE	$Z = 0 \Rightarrow PC \leftarrow \text{Oprnd}$
BRGE	$N = 0 \Rightarrow PC \leftarrow \text{Oprnd}$
BRGT	$N = 0 \wedge Z = 0 \Rightarrow PC \leftarrow \text{Oprnd}$
BRV	$V = 1 \Rightarrow PC \leftarrow \text{Oprnd}$
BRC	$C = 1 \Rightarrow PC \leftarrow \text{Oprnd}$
CALL	$SP \leftarrow SP - 2; \text{Mem}[SP] \leftarrow PC; PC \leftarrow \text{Oprnd}$
SYCALL	$Y \leftarrow \text{Mem}[\text{FFF0}]; \text{Mem}[Y - 1] \leftarrow \text{IR}\langle 0..7 \rangle; \text{Mem}[Y - 3] \leftarrow SP; \text{Mem}[Y - 5] \leftarrow PC; \text{Mem}[Y - 7] \leftarrow X;$ $\text{Mem}[Y - 9] \leftarrow A; \text{Mem}[Y - 10]\langle 4..7 \rangle \leftarrow NZVC; SP \leftarrow Y - 10; PC \leftarrow \text{Mem}[\text{FFFE}]$
LDWT	$T \leftarrow \text{Oprnd}$
LDWr	$r \leftarrow \text{Oprnd}; N \leftarrow r < 0, Z \leftarrow r = 0$
LDBr	$r\langle 8..15 \rangle \leftarrow \text{byte Oprnd}; N \leftarrow 0, Z \leftarrow r\langle 8..15 \rangle = 0$
STWr	$\text{Oprnd} \leftarrow r$
STBr	$\text{byte Oprnd} \leftarrow r\langle 8..15 \rangle$
CPWr	$Y \leftarrow r - \text{Oprnd}; N \leftarrow Y < 0, Z \leftarrow Y = 0, V \leftarrow \{overflow\}, C \leftarrow \{carry\}; N \leftarrow N \oplus V$
CPBr	$Y \leftarrow r\langle 8..15 \rangle - \text{byte Oprnd}; N \leftarrow Y < 0, Z \leftarrow Y = 0, V \leftarrow 0, C \leftarrow 0$
ADDr	$r \leftarrow r + \text{Oprnd}; N \leftarrow r < 0, Z \leftarrow r = 0, V \leftarrow \{overflow\}, C \leftarrow \{carry\}$
SUBr	$r \leftarrow r - \text{Oprnd}; N \leftarrow r < 0, Z \leftarrow r = 0, V \leftarrow \{overflow\}, C \leftarrow \{carry\}$
ANDr	$r \leftarrow r \wedge \text{Oprnd}; N \leftarrow r < 0, Z \leftarrow r = 0$
ORr	$r \leftarrow r \vee \text{Oprnd}; N \leftarrow r < 0, Z \leftarrow r = 0$
XORr	$r \leftarrow r \oplus \text{Oprnd}; N \leftarrow r < 0, Z \leftarrow r = 0$
ADDSP	$SP \leftarrow SP + \text{Oprnd}$
SUBSP	$SP \leftarrow SP - \text{Oprnd}$

Here is the memory map of the Pep/10 system. The shaded portion is ROM. Compared to the Pep/9 memory map, this map has several additional components – a disk input port at address 5555, a shutdown port at address 6666, a dispatcher at address 7777, and a unary trap handler at address 9999. Pep/9 has six machine vectors from FFF4 to FFFE. Because of the new components, Pep/10 has ten machine vectors at addresses FFEC to FFFE.



Here is the data section of the Pep/10 CPU. Compared to the Pep/9 data section, Pep/10 has two additional components – a shadow carry bit, denoted S in the figure below, and an additional multiplexer with its associated control line CSMux. The shadow carry bit is not visible at the ISA level and is used for internal address calculations in the microcode. This design solves a major headache present in Pep/9, which requires the saving and restoration of the C bit when an internal address addition would wipe it out. A step towards a more realistic model is the requirement of three consecutive MemRead/MemWrite assertions for memory access as opposed to two with Pep/9.



Here is the data section of the Pep/10 CPU with the two-byte data bus. The fifth edition of *Computer Systems* drops the discussion of the MAR Incrementer in favor of a more extensive discussion of increasing the data bus width to improve performance. The material is improved by incorporating it into the Pep9CPU software. Students can toggle between the two models, with and without the wider data bus, test their solutions with the software, and use the UnitPre and UnitPost tests in the Help system.

